

! !

3 - LS

,

Table of content

- 1.
- 2.
- 3.
- 4.
- 5.

- : **3**
 - : , , ,
-

- ! , ,
 - , , !
 - !
-
-

1. .
2. .
3. .
4. .

-

-

```
import pandas as pd
import numpy as np
import matplotlib as mpl

mpl.rcParams['font.family'] = 'Malgun Gothic'
camera = pd.read_csv('../team2/data/camera.csv')
accident = pd.read_csv('../team2/data/accident.csv', encoding='cp949')
```

-

```
print(camera.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27403 entries, 0 to 27402
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0           27403 non-null  object
1           27403 non-null  object
2           27403 non-null  object
```

```

3          27403 non-null  object
4          7964 non-null  object
5          27403 non-null  object
6          27403 non-null  int64
7          11385 non-null object
8          24525 non-null object
9          27403 non-null  float64
10         27403 non-null  float64
11         27403 non-null  object
12         27403 non-null  int64
13         27403 non-null  int64
14         1123 non-null   float64
15         954 non-null   float64
16         27403 non-null  int64
17         27403 non-null  int64
18         27403 non-null  object
19         27403 non-null  object
20         27403 non-null  object
dtypes: float64(4), int64(5), object(12)
memory usage: 4.4+ MB
None

```

•

```
print(accident.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 230 entries, 0 to 229
Data columns (total 7 columns):
 #   Column  Non-Null Count  Dtype
---  -
0          230 non-null  object
1          230 non-null  object
2          230 non-null  int64
3          230 non-null  int64
4          230 non-null  int64
5          230 non-null  int64
6          230 non-null  int64
dtypes: int64(5), object(2)
memory usage: 12.7+ KB
None

```

[illegible]

```
#
theeshold = accident[''].quantile(0.75)
accident[''] = accident[''].apply(lambda x: '' if x >= theeshold else '')

#
road_type = camera.groupby('')[''].agg(lambda x: x.value_counts().index[0])
road_type = road_type.rename(columns={'': ''})
```

```
#
acc_road = pd.merge(road_type, accident, on=' ')

# vs
ct = pd.crosstab(acc_road[' '], acc_road[' '])

ct.loc[' '] = ct.loc[' '] + ct.loc[' ']
ct = ct.drop([' '])
```

```
#
ct.style.set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#f4d9c6'),
                                ('color', '#333'),
                                ('text-align', 'center'),
                                ('font-family', 'Jua'),
                                ('font-size', '1.1em')]}],
    {'selector': 'td', 'props': [('text-align', 'center'),
                                ('font-family', 'Jua'),
                                ('padding', '10px')]}],
    {'selector': '', 'props': [('border', '1px solid #ddd'),
                                ('border-collapse', 'collapse')]}]
).set_caption(" ") \
.set_properties(**{
    'background-color': '#ffdf9',
    'border-color': '#eee',
    'border-style': 'solid',
    'border-width': '1px'
})
```

Table 1:

13	37
25	74
6	36
13	12

-
- 1 :

```
#
from scipy.stats import chi2_contingency
import matplotlib.pyplot as plt
chi2, p, dof, exp = chi2_contingency(ct)

print(chi2, p, dof)
```

11.678710326416384 0.008568853463999705 3

- (Chi²): **11.68** (df): **3** (p-value): **0.008**
!

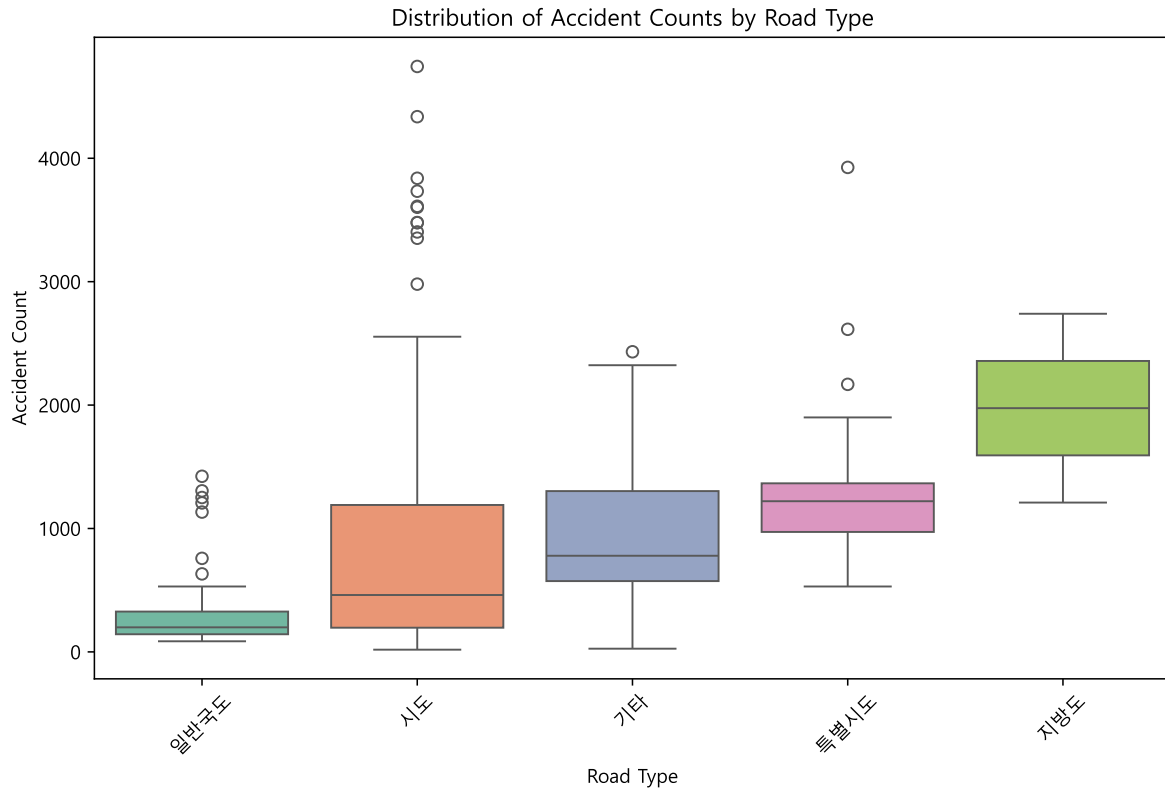
- [!]

-

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(10,6))
sns.boxplot(x=' ', y=' ', data=acc_road, palette='Set2')
plt.xticks(rotation=45)
plt.title('Distribution of Accident Counts by Road Type')
plt.xlabel('Road Type')
plt.ylabel('Accident Count')
plt.show()
```

C:\Users\USER\AppData\Local\Temp\ipykernel_11132\891112310.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign




```

ct_norm = ct.div(ct.sum(axis=1), axis=0) #
# long-format
ct_long = ct_norm.reset_index().melt(id_vars=' ', var_name=' ', value_name=' ')

plt.figure(figsize=(10, 6))
sns.pointplot(data=ct_long, x=' ', y=' ', hue=' ', dodge=0.3, markers='o', linestyle='--')
plt.xticks(rotation=45)
plt.title(' / ')
plt.ylabel(' ')
plt.xlabel(' ')
plt.tight_layout()
plt.show()

```



```
# accident ( )
threshold = accident[' '].quantile(0.75)
accident[' '] = accident[' '].apply(lambda x: ' ' if x >= threshold else ' ')

# camera accident ( )
all_road = pd.merge(camera[[' ', ' ']], accident[[' ', ' ', ' ']], on=' ')


#
ct_all = pd.crosstab(all_road[' '], all_road[' '])
```

```
#
ct_all_norm = ct_all.div(ct_all.sum(axis=1), axis=0)
ct_all_long = ct_all_norm.reset_index().melt(id_vars=' ', var_name=' ', value_name=' ')

plt.figure(figsize=(10, 6))
sns.pointplot(data=ct_all_long, x=' ', y=' ', hue=' ', dodge=0.3, markers='o', linestyle=' ')
plt.xticks(rotation=45)
```



```
plt.ylabel(' ')
plt.xlabel(' ')
plt.title(' / (Point Plot)')
plt.tight_layout()
plt.show()
```



- Shapiro-Wilk

```
from scipy.stats import shapiro
def check_normality_by_road_type(all_road):
    #
    results = []

    #
    for name, group in all_road.groupby(' '):
        stat, p = shapiro(group[' '])
        result = ' ' if p < 0.05 else ' '
        results.append([name, p, result])
```

```
#
df_results = pd.DataFrame(results, columns=['', 'p-value', ''])

#
return df_results
```

```
df_normality = check_normality_by_road_type(all_road)
df_normality
```

```
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 8801.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 26516.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 53789.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 216655.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 77536.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 57216.
C:\Users\USER\.conda\envs\ls_pyun\Lib\site-packages\scipy\stats\_axis_nan_policy.py:586: UserWarning:
scipy.stats.shapiro: For N > 5000, computed p-value may not be accurate. Current N is 53306.
```

	p-value
0	6.372202e-78
1	4.453160e-29

	p-value
2	8.578405e-119
3	3.224323e-109
4	1.458738e-146
5	6.107652e-135
6	2.914311e-122
7	4.290432e-124

-
- Kruskal Wallis H

```
from scipy.stats import kruskal

groups = [g[' '].values for _, g in all_road.groupby(' ')]
stat, p = kruskal(*groups)
print("Kruskal-Wallis H :", p)
```

Kruskal-Wallis H : 0.0

-
- Kruskal Wallis H

[!]

-
- 2 :

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

#
plt.rcParams['font.family'] = 'Malgun Gothic'
plt.rcParams['axes.unicode_minus'] = False

accident_df = pd.read_csv('../team2/data/accident.csv', encoding='cp949')

# unique
unique_values = accident_df[' '].unique()
```

•

```
#
accident_avg = accident_df.groupby(' ')[ ' '].mean().reset_index()
```

•

```
ac_avg = accident_avg.sort_values(by=' ', ascending=False)
ac_avg.head(7)
```


14	1975.000000
1	1714.967742
6	1479.600000
4	1394.400000
8	1352.440000
5	1208.888889
9	1166.000000

•

```
#
plt.figure(figsize=(12, 6))
sns.boxplot(data=accident_df, x=' ', y=' ', palette='Set3')
plt.title(' ')
plt.xlabel(' ')
plt.ylabel(' ')
plt.xticks(rotation=45)
plt.tight_layout()
plt.grid(True)
plt.show()
```

C:\Users\USER\AppData\Local\Temp\ipykernel_11132\4283594563.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign



•

[!]

!

• ANOVA

```
# : ANOVA
import statsmodels.api as sm
from statsmodels.formula.api import ols

model = ols(' ~ C( )', data=accident_df).fit()

anova_results = sm.stats.anova_lm(model, typ=2)
```

H0:

.

H1:

.

```

from IPython.display import HTML

anova_results = pd.DataFrame(anova_results)

styled = anova_results.style.set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#f4d9c6'),
                                  ('color', '#333'),
                                  ('text-align', 'center'),
                                  ('font-family', 'Jua'),
                                  ('font-size', '1.1em')]}],
    {'selector': 'td', 'props': [('text-align', 'center'),
                                  ('font-family', 'Jua'),
                                  ('padding', '10px')]}],
    {'selector': '', 'props': [('border', '1px solid #ddd'),
                                ('border-collapse', 'collapse')]}]
).set_caption(" ANOVA ")

HTML(styled.to_html())

```

Table 4: ANOVA

	sum_sq	df	F	PR(>F)
C()	55213437.753574	16.000000	5.600300	0.000000
Residual	131248122.611644	213.000000	nan	nan

-
- ANOVA

[!]

-
- Kruskal Wallis H

```

from scipy.stats import kruskal

grouped_values = [group[' '].values for _, group in accident_df.groupby(' ')]
stat, p = kruskal(*grouped_values)
stat, p

```

(89.2757873932235, 3.4006688531866835e-12)

-
- Kruskal Wallis H


[!] !

- : Dunn's test

```
import scikit_posthocs as sp_post

posthoc = sp_post.posthoc_dunn(accident, val_col=' ', group_col=' ', p_adjust='bonferroni')
# Bonferroni

#
plt.figure(figsize=(10, 7))
sns.heatmap(posthoc,
            annot=True,
            fmt=".3f",
            cmap="coolwarm_r",
            cbar_kws={'label': 'p-value'},
            linewidths=0.5,
            square=True)
plt.title("Dunn's Test: (p-value heatmap)")
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```



- 3 :

```
accident_df = pd.read_csv('../team2/data/accident.csv', encoding='cp949')
c_df = pd.read_csv('../team2/data/camera.csv')

sido_map = {
    '경기도': '1',
    '경상남도': '2',
    '경상북도': '3',
    '광주광역시': '4',
    '대전광역시': '5',
    '충청남도': '6',
    '충청북도': '7',
    '전라남도': '8',
    '전라북도': '9',
    '제주특별자치도': '10',
    '서울특별시': '11',
    '부산광역시': '12',
    '울산광역시': '13',
    '인천광역시': '14',
    '강원특별자치도': '15',
    '세종특별자치시': '16'
}
```



```
agg_cam = camera.groupby(' ',as_index=False)[' '].count()
agg_cam.columns = ['SIDO_NM',' ']
```

```
# plotly
fig = px.choropleth_mapbox(
agg_cam,
geojson=geojson_data,
locations="SIDO_NM",
featureidkey="properties.SIDO_NM",
color=" ",
color_continuous_scale="Blues",
mapbox_style="carto-positron",
center={"lat": 37.5665, "lon": 126.9780},
zoom=5,
opacity=1,
title=" ",
)
fig.update_layout(margin={"r":0,"t":30,"l":0,"b":0})
fig.show()
```

```
# 0
camera_count = c_df.groupby([' ', ' ']).size().reset_index(name=' ')
camera_count.rename(columns={' ': ' '}, inplace=True)

#
camera_merged_df = pd.merge(camera_count, accident_df.groupby([' ', ' '])[' '].sum().reset_index(),
                             on=[' ', ' '])

# ( / / )
q1 = camera_merged_df[' '].quantile(0.25)
q3 = camera_merged_df[' '].quantile(0.75)


def classify_group(x):
    if x <= q1:
        return ' '
    elif x >= q3:
        return ' '
    else:
        return ' '
```

```
camera_merged_df[' '] = camera_merged_df[' '].apply(classify_group)
```

```
# Boxplot :
plt.figure(figsize=(8,5))
sns.boxplot(x=' ', y=' ', data=camera_merged_df, palette='pastel')
plt.title(' ')
plt.xlabel(' ')
plt.ylabel(' ')
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\USER\AppData\Local\Temp\ipykernel_11132\1637060997.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign



-
- ANOVA

```
from scipy.stats import f_oneway

group_ = camera_merged_df[camera_merged_df[' '] == ' '][' ']
group_ = camera_merged_df[camera_merged_df[' '] == ' '][' ']
group_ = camera_merged_df[camera_merged_df[' '] == ' '][' ']

f_stat, p_value = f_oneway(group_ , group_ , group_ )
# p-value < 0.05,
```

```
from IPython.display import HTML

anova_result_df = pd.DataFrame({
    'F- ': [f_stat],
    'p- ': [p_value]
})

styled = anova_result_df.style.set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#f4d9c6'),
                                  ('color', '#333'),
                                  ('text-align', 'center'),
                                  ('font-family', 'Jua'),
                                  ('font-size', '1.1em')]}],
    {'selector': 'td', 'props': [('text-align', 'center'),
                                  ('font-family', 'Jua'),
                                  ('padding', '10px')]}],
    {'selector': '', 'props': [('border', '1px solid #ddd'),
                                ('border-collapse', 'collapse')]}]
).set_caption(" ANOVA ")

#
HTML(styled.to_html())
```

Table 5: ANOVA

	F-	p-
0	43.666575	0.000000

H0: $\mu_1 = \mu_2 = \mu_3 = \mu_4$ H1: $\mu_1 \neq \mu_2 \neq \mu_3 \neq \mu_4$

- Shapiro-Wilk / Levene

```
results = []
# shapiro:
from scipy.stats import shapiro

for name, group in zip(['1', '2', '3', '4'], [group_1, group_2, group_3, group_4]):
    stat, p = shapiro(group)
    norm_check = 'OK' if p > 0.05 else 'X'
    results.append([name, round(p, 4), norm_check])

# Levene:
from scipy.stats import levene

stat, p = levene(group_1, group_2, group_3, group_4)
equal_check = 'OK' if p > 0.05 else 'X'
results.append(['', round(p, 4), equal_check])

df_results = pd.DataFrame(results, columns=['', 'p-value', ''])

df_results
```

	p-value	
0	0.0000	X
1	0.0000	X
2	0.0003	X
3	0.0000	X

- Kruskal Wallis H

```
from scipy.stats import kruskal

h_stat, p_value = kruskal(group_ , group_ , group_ )
print(f"Kruskal-Wallis H      : {h_stat:.3f}, p-value: {p_value:.4f}")
```

Kruskal-Wallis H : 71.014, p-value: 0.0000

-
- Kruskal Wallis H

[!] !

- : Dunn's test

```
import scikit_posthocs as sp

dunn_result = sp.posthoc_dunn(camera_merged_df, val_col=' ', group_col=' ', p_adjust='bonfe

#
pairs = []
for i, group_a in enumerate(dunn_result.index):
    for j, group_b in enumerate(dunn_result.columns):
        if i < j: # ( )
            pairs.append({
                'A': group_a,
                'B': group_b,
                'p-value': dunn_result.iloc[i, j]
            })

#
dunn_pairs_df = pd.DataFrame(pairs)

dunn_pairs_df
```

	A	B	p-value
0			1.598669e-07
1			1.399293e-16

A	B	p-value
2		6.733709e-05

- 4 :

```
#             True,      False
c_df['      '] = c_df['      '].notna()

#
roadline_status = c_df.groupby([' ', ' '])['      '] \
                    .agg(lambda x: x.value_counts().index[0]) \
                    .reset_index(name='      ')

roadline_status.rename(columns={'      ': '      '}, inplace=True)


accident_with_roadline = pd.merge(
    accident_df[[' ', ' ', ' ']],
    roadline_status,
    on=[' ', ' '])
```

```
accident_with_roadline['      '] = accident_with_roadline['      '].map({True: ' ', False: ' '})

plt.figure(figsize=(6, 4))
sns.boxplot(data=accident_with_roadline, x='      ', y='      ', palette='Set3')
plt.title('      ')
plt.xlabel('      ')
plt.ylabel('      ')
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\USER\AppData\Local\Temp\ipykernel_11132\2032049726.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign



- 2 T

```
from scipy.stats import shapiro, ttest_ind, mannwhitneyu

group_with = accident_with_roadline[accident_with_roadline['   '] == ' ']['   ']
group_without = accident_with_roadline[accident_with_roadline['   '] == ' ']['   ']

t_stat, p_value = ttest_ind(group_with, group_without, equal_var=False)

results = t_stat, p_value
```

- 2 T


```

        = ' ' if p_value < 0.05 else ' '

t_stat, p_value = ttest_ind(group_with, group_without, equal_var=False)

results = [[' ', f"{p_value:.3f}", ' ']]

# DataFrame
df_results = pd.DataFrame(results, columns=[' ', 'p_value', ' '])

df_results

```

	p_value
0	0.000

-
- Shapiro-Wilk / Mann-Whitney U

```

from scipy.stats import shapiro, ttest_ind, mannwhitneyu

stat, p = shapiro(group_with)
out_stat, out_p = shapiro(group_without)

stat, p = mannwhitneyu(group_with, group_without)

```

-
- Shapiro-Wilk / Mann-Whitney U

```

#
results = []
stat1, p1 = shapiro(group_with)
result1 = ' ' if p1 > 0.05 else ' '
results.append([' ', round(p1, 4), result1])

stat2, p2 = shapiro(group_without)
result2 = ' ' if p2 > 0.05 else ' '
results.append([' ', round(p2, 4), result2])

# (Mann-Whitney U)

```

```

mw_stat, mw_p = mannwhitneyu(group_with, group_without)

# DataFrame
results.append(['Mann-Whitney U ', round(mw_p, 4), ' ' if mw_p < 0.05 else ' '])

df_results = pd.DataFrame(results, columns=[' / ', 'p-value', ' '])

styled_df = df_results.style.set_table_attributes('class="styled-table"')

styled_df

```

Table 9

	/	p-value
0		0.000000
1		0.000000
2	Mann-Whitney U	0.000000

- Levene / Brunner Munzel

```

#
from scipy.stats import levene
stat, p = levene(group_with, group_without)

from scipy.stats.mstats import brunnermunzel
stat, pvalue = brunnermunzel(group_with, group_without, alternative='two-sided')

```

- Levene / Brunner Munzel

```

#
results = []
stat, p = levene(group_with, group_without)
result1 = ' ' if p1 > 0.05 else ' '
results.append(['Levene', round(p1, 4), result1])

stat, pvalue = brunnermunzel(group_with, group_without, alternative='two-sided')

```

```

result2 = ' ' if p2 > 0.05 else ' X'
results.append(['BM', round(p2, 4), result2])

df_results = pd.DataFrame(results, columns=[' ', 'p-value', ' '])

styled_df = df_results.style.set_table_attributes('class="styled-table"')

styled_df

```

Table 10

		p-value	
0	Levene	0.000000	
1	BM	0.000000	X

- Brunner Munzel

[!]

-

```

camera[' '] = camera[' '].notna()
camera.rename(columns={' ': ' '}, inplace=True)

camera_count = camera.groupby([' ', ' ']).size().reset_index(name=' ')
camera_merged_df = pd.merge(camera_count, accident.groupby([' ', ' '])[' '].sum().reset_index(), on=[' ', ' '])

merged_df = pd.merge(camera_merged_df, camera[[' ', ' ', ' ', ' ']], on=[' ', ' '])

summary_df = (
    merged_df.groupby([' ', ' '])
    .agg({
        ' ': 'first',
        ' ': 'first',
        ' ': 'mean' # True
    })
    .reset_index()
)

```

```
)

#
summary_df = summary_df[[' ', ' ', ' ', ' ']]
summary_df.rename(columns={' ': ' '}, inplace=True)
```

•

```
pop_df = pd.read_csv('../team2/data/ _ _ .csv')

pop_df = pop_df.rename(columns={pop_df.columns[0]: ' ', pop_df.columns[1]: ' '})
pop_df = pop_df[~pop_df[' '].str.contains(' | ')] #
pop_df = pop_df.dropna() # NaN
pop_df[' '] = pop_df[' '].astype(int)

final_df = pd.merge(summary_df, pop_df, on=' ', how='left')
```

•

```
from statsmodels.formula.api import ols

model = ols(' ~ + + ', final_df).fit()
```

1. y :
2. X : , , /

```
print(model.summary())
```

```

OLS Regression Results
=====
Dep. Variable:                R-squared:                0.695
Model:                        OLS    Adj. R-squared:                0.692
Method:                        Least Squares    F-statistic:                244.1
Date:                        Thu, 17 Apr 2025    Prob (F-statistic):        1.38e-82
```

Time: 20:00:30 Log-Likelihood: -2449.5
 No. Observations: 326 AIC: 4907.
 Df Residuals: 322 BIC: 4922.
 Df Model: 3
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	110.4797	55.679	1.984	0.048	0.939	220.021
	3.2748	0.384	8.517	0.000	2.518	4.031
	-305.3707	109.047	-2.800	0.005	-519.905	-90.836
	0.0022	0.000	12.598	0.000	0.002	0.003
Omnibus:	71.350		Durbin-Watson:	1.188		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	196.981		
Skew:	1.004		Prob(JB):	1.68e-43		
Kurtosis:	6.236		Cond. No.	1.39e+06		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 1.39e+06. This might indicate that there are strong multicollinearity or other numerical problems.

1. Adj. R-squared: 0.699 !
2. Prob (F-statistic): 2.44e-87 !

• OLS

```
coef = model.params
p_values = model.pvalues
conf_int = model.conf_int()

# ( )
variables = ['Intercept', ' ', ' ', ' ']

#
summary_table = pd.DataFrame({
    ' ': variables,
    'coef': [coef[var] for var in variables],
    'p': [p_values[var] for var in variables],
```

```

    '0.025': [conf_int.loc[var][0] for var in variables],
    '0.975': [conf_int.loc[var][1] for var in variables]
})

r2 = model.rsquared
adj_r2 = model.rsquared_adj

#
display(summary_table.round(3))

```

		coef	p	0.025	0.975
0	Intercept	110.480	0.048	0.939	220.021
1		3.275	0.000	2.518	4.031
2		-305.371	0.005	-519.905	-90.836
3		0.002	0.000	0.002	0.003

1. p-value 5%
2. “ ”

```

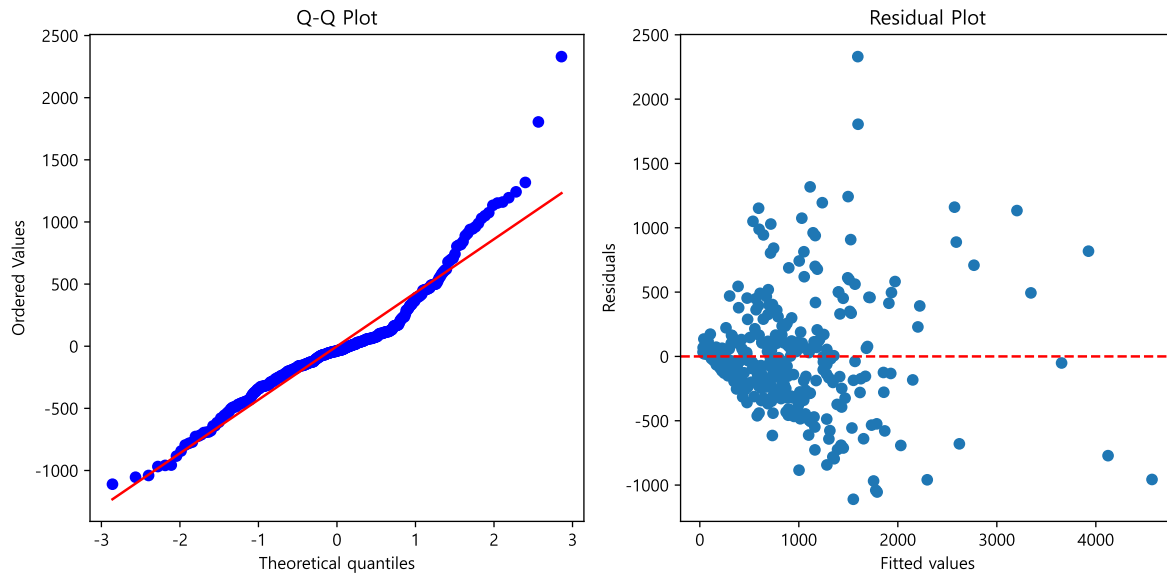
import matplotlib.pyplot as plt
import scipy.stats as stats
#| echo: true
#| code-fold: true
residuals = model.resid
#      (1 2)
fig, axes = plt.subplots(1, 2, figsize=(10, 5))

# Q-Q plot
stats.probplot(residuals, dist="norm", plot=axes[0])
axes[0].set_title("Q-Q Plot")

# Residual plot
axes[1].scatter(model.fittedvalues, residuals)
axes[1].axhline(0, color='red', linestyle='--')
axes[1].set_xlabel('Fitted values')
axes[1].set_ylabel('Residuals')
axes[1].set_title('Residual Plot')

```

```
#
plt.tight_layout()
plt.show()
```



```
import scipy.stats as sp
W, p = sp.shapiro(model.resid)
```

•

```
results = []
W, p = shapiro(model.resid)
result1 = ' ' if p > 0.05 else ' X'
results.append(['shapiro', round(p, 4), result1])

df_results = pd.DataFrame(results, columns=[' ', 'p-value', ' '])

styled_df = df_results.style.set_table_attributes('class="styled-table"')
styled_df
```

Table 12

p-value			
0	shapiro	0.000000	X

, Durbin-Watson !

.

•

, , , , .

•

1. .
2. .
3. .

!